


# blocksize


2025-04-12

tysons



*Freedom of Speech, 1943*

1. Why it's worth thinking about
2. Measurements and demythologizing



# My goals

1. Deliver on bitcoin's promise of value transmission/storage without intermediaries
  - a. L1 “checking acct volume” for most people (1 txn/month)
2. Avoid (de facto) state capture

# My goals

1. Deliver on bitcoin's promise of value transmission/storage without intermediaries
  - a. L1 “checking acct volume” for most people (1 txn/month)
2. Avoid (de facto) state capture

**This is not about payments. This is about store of value.**

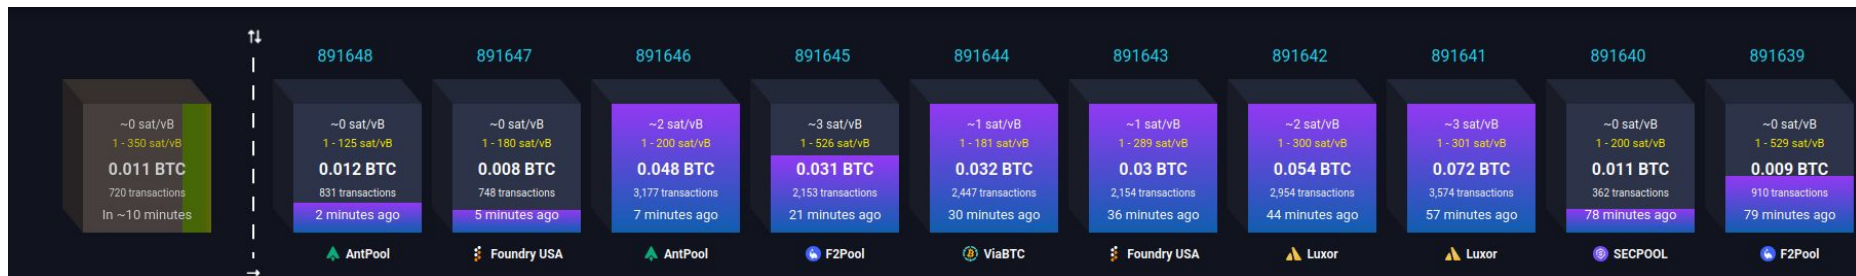
# My goals

1. Deliver on bitcoin's promise of value transmission/storage without intermediaries
  - a. L1 “checking acct volume” for most people (1 txn/month)
2. Avoid (de facto) state capture

**This is not about payments. This is about store of value.**


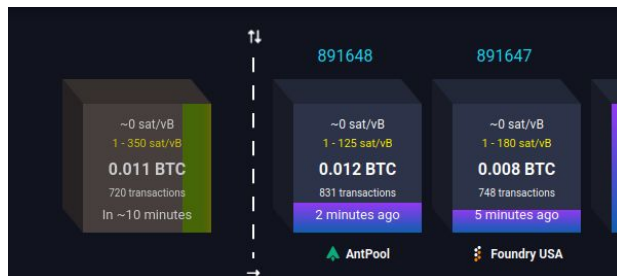
**I think 3 is mostly achievable with tech we have (or know about) today**

# Why care?







# Why care?



# Why care?

“well, I can still  
transact”




# Finance is discontinuous

## Mammoth Outflows

Credit Suisse sees historic levels of clients pulling their money

■ Net New Money

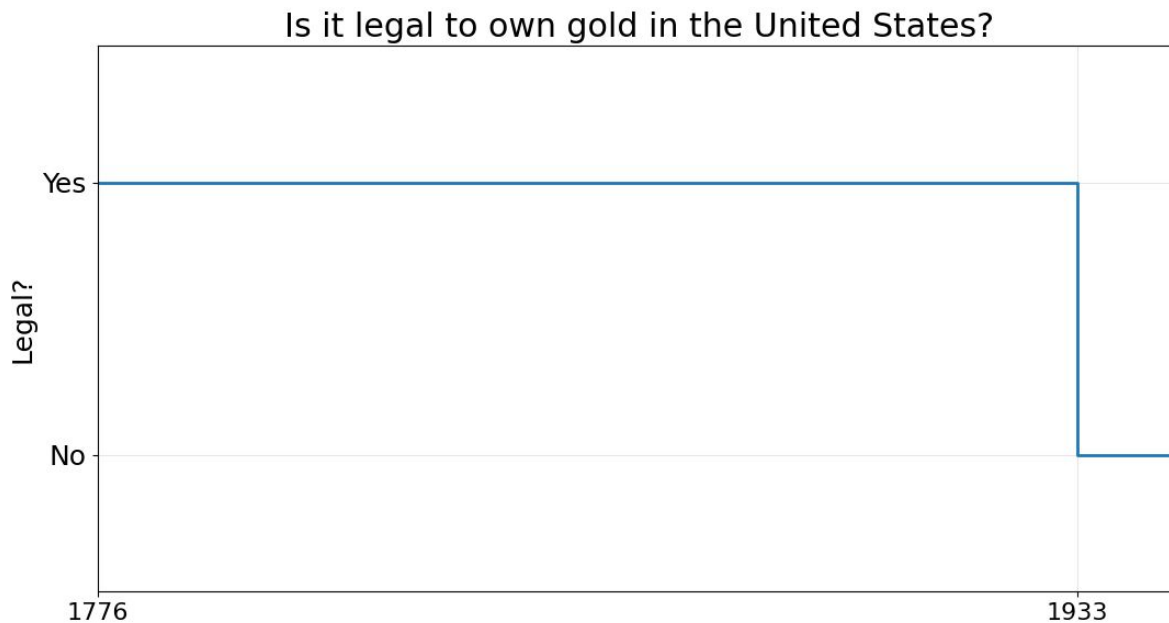


Source: Bloomberg Intelligence and company financial statements.

**Bloomberg**

regulation

# ~~Finance~~ is discontinuous



“On the afternoon of the Wednesday before Thanksgiving, something unexpected will happen to the turkey. It will incur a revision of belief.”

Nassim Taleb (ardent bitcoin supporter)

not your keys,

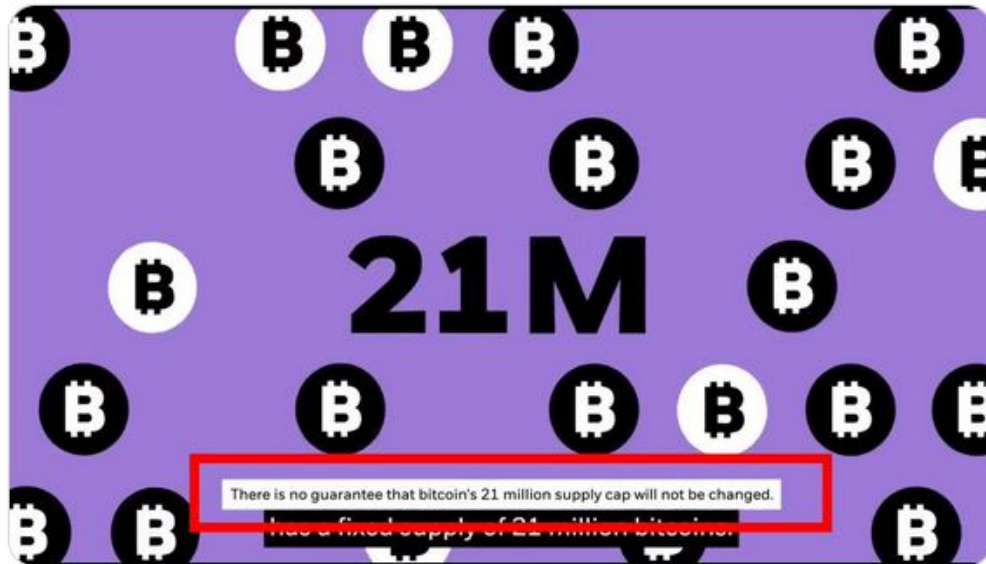
not your ~~coins~~ consensus



jamesob  
@jamesob



Bitcoin™!



5:50 PM · Dec 18, 2024 · 136.9K Views

View post engagements



14



34



156



11



all of the sudden  
everyone wants to be a self-custodian

now what?



# chain limitations on exit

```
[10]: # Per https://contenthub-static.crypto.com/wp\_media/2024/08/Crypto.com-Crypto-Market-Sizing-H1-2024-1.pdf
      est_bitcoin_users = 314_000_000

      # Assume that 5% of Bitcoin holders are in self-custody (probably generous)
      est_self_custody_users = 0.05 * est_bitcoin_users
      est_exch_users = est_bitcoin_users - est_self_custody_users

      print(f"there are {est_exch_users / 1e6:_}M users holding coins on exchange")

      there are 298.3M users holding coins on exchange
```

# chain limitations on exit

```
[2]: vb_per_block = 1_000_000

blocks_per_month = 6 * 24 * 30

txsize_2in2out_vb = 208.5 # per https://bitcoinops.org/en/tools/calc-size/

txns_per_block = int(vb_per_block / txsize_2in2out_vb)

txns_per_month = txns_per_block * blocks_per_month

print(f"capacity: {txns_per_month:,} (2-in-2-out) txns per month")
```

capacity: 20,718,720 (2-in-2-out) txns per month

# chain limitations on exit

```
[10]: # Per https://contenthub-static.crypto.com/wp\_media/2024/08/Crypto.com-Crypto-Market-Sizing-H1-2024-1.pdf
est_bitcoin_users = 314_000_000

# Assume that 5% of Bitcoin holders are in self-custody (probably generous)
est_self_custody_users = 0.05 * est_bitcoin_users
est_exch_users = est_bitcoin_users - est_self_custody_users

print(f"there are {est_exch_users / 1e6:_}M users holding coins on exchange")

there are 298.3M users holding coins on exchange
```

```
[12]: custodial_move_months = est_exch_users / txns_per_month
print(f"it would take {custodial_move_months:.2f} "
      f"months of uninterrupted use chain for the {est_exch_users / 1e6:_}M custodial users to withdraw")

it would take 14.40 months of uninterrupted use chain for the 298.3M custodial users to withdraw
```

1. US gov (or G7 countries) 6102 an impl. w/ forks
2. “Exit window” given
3. Blackrock has fork choice in their S1, dumps other side
4. Profit-seeking (and regulated) miners must follow
5. “Real” bitcoin PoW cooked

1. US gov (or G7 countries) 6102 an impl. w/ forks
2. “Exit window” given
3. Blackrock has fork choice in their S1, dumps other side
4. Profit-seeking (and regulated) miners must follow
5. “Real” bitcoin PoW cooked



# deployed L2s don't solve this

```
[14]: # A 1-in-2-out LN channel. Bare minimum, probably larger on average.  
txsize_ln_open_vb = 154 # e.g. https://mempool.space/tx/f9b47a6c2137c0adf0cbb44155b46b27b690018fc16eb282abb85e791ca90715  
  
time_to_US_channels = (est_exch_users * txsize_ln_open_vb) / (vb_per_block * blocks_per_month)  
  
print(f"it would take {time_to_US_channels:.2f} months of uninterrupted chain space "  
      f"to open an LN channel for each custodial user")
```

it would take 10.63 months of uninterrupted chain space to open an LN channel for each custodial user

# best-case response

- some kind of batched withdrawal to “other” L2s
  - Ark, Lightning, cashu, fedimint
  - some TBD shared UTXO scheme
- you don't have real bitcoin, but you escaped gov?
- exchanges need preexisting batch connections

# best-case response

- some kind of batched withdrawal to “other” L2s
  - Ark, Lightning, cashu, fedimint
  - some TBD shared UTXO scheme
- you don't have real bitcoin, but you escaped gov?
- exchanges need preexisting batch connections
- **without L1 access, you can always get rugged**




# Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto  
satoshin@gmx.com  
www.bitcoin.org

**Abstract.** A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network.

a more optimistic case

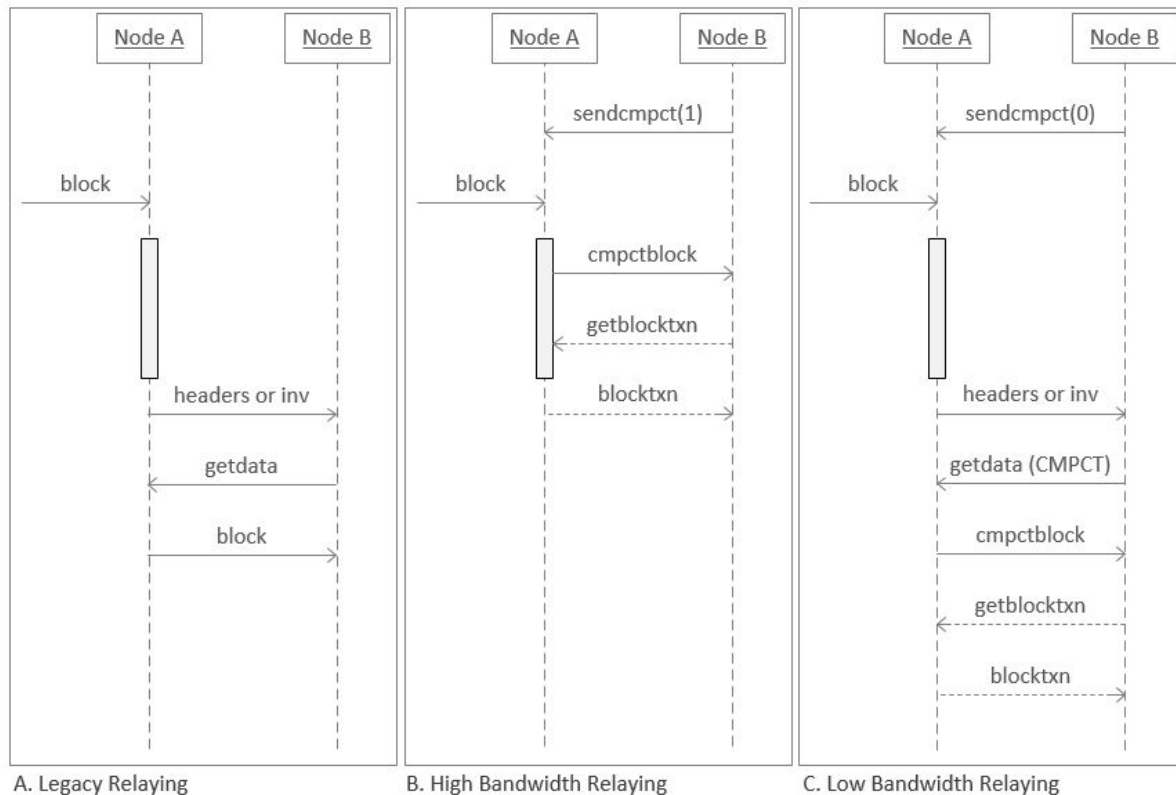


objections

# miner centralization

block propagation

# compact blocks



# miner centralization

block propagation

good hardware, internet connection is basically the last thing  
limiting pools

# mempool diffs

“too much in the mempool; upon new block I’ll have to retrieve too much”

# mempool diffs

- Weak blocks fix this
  - Like pool shares for the network
  - PoW backed mempool syncing
  - Miner-incentive compatible
- P2P relay was always sort of doomed anyway
  - Relies on bandwidth donation
  - Useful policy always lags
- Would need this anyway with high fee rates



# Second Look at Weak Blocks

■ Implementation



instagibbs

2  Apr 2024

[Weak blocks](#) <sup>28</sup>, or “near blocks” are not a new idea.

In short, have miners propagate what amounts to mining shares over the p2p network, which allows PoW-backed sharing of data.

[Historical discussions](#) <sup>8</sup> of weak blocks centered around the [blocksize and scaling debate](#) <sup>4</sup>, which means there was intense focus on reducing the marginal bytes sent per weak block to aid “gigameg blocks”. There was seemingly a lot of focus on creating DAGs, extra-consensus chains, and similar mechanisms for increasing the blocksize safely.

Almost 10 years have passed, communities have split, basically everyone is a small blocker of some kind. Ignoring blocksize increases as a motivation, *is there value in reconsidering this type of proposal?*

# miner centralization

selfish mining

- mostly about share of hash rather than blocksize
- worth thinking about

# fee revenue for miners

“miners will need high fees and  
full blocks to get paid”

# fee revenue for miners

Assuming a price of \$250,000/BTC (i.e. \$0.0025/sat)

Fee subsidy of 1.5625 BTC

- |                                 |                                      |
|---------------------------------|--------------------------------------|
| - 4MvB block                    | - 1 sat/vB feerate                   |
| - feerate required: 156 sats/vB | - full blocksize 156.25 MvB required |
| - total payments: 4,796         | - total payments: 749,400            |
| - cost per payment: \$81.45     | - cost per payment: \$0.52           |

Fee subsidy of 3.125 BTC

- |                                 |                                     |
|---------------------------------|-------------------------------------|
| - 4MvB block                    | - 1 sat/vB feerate                  |
| - feerate required: 312 sats/vB | - full blocksize 312.5 MvB required |
| - total payments: 4,796         | - total payments: 1,498,800         |
| - cost per payment: \$162.89    | - cost per payment: \$0.52          |

Fee subsidy of 6.25 BTC

- |                                 |                                     |
|---------------------------------|-------------------------------------|
| - 4MvB block                    | - 1 sat/vB feerate                  |
| - feerate required: 625 sats/vB | - full blocksize 625.0 MvB required |
| - total payments: 4,796         | - total payments: 2,997,601         |
| - cost per payment: \$325.78    | - cost per payment: \$0.52          |



# makes IBD harder

“how will I sync?”

# ~~makes IBD harder~~

- easy problem to solve
- assumeutxo is already sufficient
  - to tip in under ~2 hours on most computers
  - assumevalid already in use by default
- maybe someday
  - swiftsync
  - zerosync

# makes tip maintenance harder

`“how will my node stay current?”`



# brief measurement digression

- Measurements performed on “pretty good” consumer desktop
- Directionally true to anything with SSD, >32GB RAM

# let's make it interesting

```
[6]: ppl = 7_000_000_000
    target_txsperpersonmonth = 2
    target_txspermonth = ppl * target_txsperpersonmonth

    txs_per_block = int(target_txspermonth / blocks_per_month)

    mvb_per_block = (txs_per_block * txsize_2in2out_vb) // (1000 ** 2)
    sig_checks_per_block = txs_per_block * 2

    print(f"MvB per block: {mvb_per_block}")
    print(f"sig checks per block: {sig_checks_per_block:_}")
```

```
MvB per block: 675.0
sig checks per block: 6_481_480
```

# checking sigs

File: secpbench.cpp

```
1  #include <iostream>
2  #include <format>
3  #include <cstring>
4  #include <vector>
5  #include <thread>
6  #include <random>
7  #include <chrono>
8  #include <secp256k1.h>
9
10 constexpr size_t NUM_KEYS = 6'500'000;
11 constexpr size_t NUM_THREADS = 30;
12
13 void generate_keys(std::vector<std::vector<unsigned char>>& keys, size_t start, size_t end) {
14     std::random_device rd;
15     std::mt19937_64 gen(rd()); // 64-bit generator
16     std::uniform_int_distribution<uint64_t> dist(0, UINT64_MAX);
17
18     for (size_t i = start; i < end; ++i) {
19         std::vector<unsigned char> key(32);
20         for (size_t j = 0; j < 4; ++j) { // Generate 4 blocks of 64 bits (8 bytes each)
21             uint64_t rand_block = dist(gen);
22             std::memcpy(key.data() + j * 8, &rand_block, 8);
23         }
24         keys[i] = std::move(key);
25         if (i % 10'000 == 0) {
26             std::cout << "Generated " << i << " keys" << std::endl;
27         }
28     }
29 }
```

# checking sigs

```
30
31 void benchmark_sign(
32     const std::vector<std::vector<unsigned char>>& keys, size_t thread_id, size_t& sign_count) {
33     secp256k1_context* ctx = secp256k1_context_create(SECP256K1_CONTEXT_SIGN);
34     unsigned char msg[32] = {0};
35     unsigned char sig[64];
36     size_t count = 0;
37     size_t keys_per_thread = NUM_KEYS / NUM_THREADS;
38     size_t start = thread_id * keys_per_thread;
39     size_t end = (thread_id + 1) * keys_per_thread;
40
41     for (size_t i = start; i < end; ++i) {
42         secp256k1_ecdsa_signature signature;
43         if (secp256k1_ecdsa_sign(ctx, &signature, msg, keys[i].data(), nullptr, nullptr)) {
44             ++count;
45         }
46         if (count % 10'000 == 0) {
47             std::cout << std::format("signer thread {} signed {}\n", thread_id, count);
48         }
49     }
50
51     secp256k1_context_destroy(ctx);
52     sign_count = count;
53 }
54
55 keys[i] = std::move(key);
56 if (i % 10'000 == 0) {
57     std::cout << "Generated " << i << " keys" << std::endl;
58 }
59 }
```

checking sigs



# checking sigs

```
12:55:42 james@fido james/tmp % clang++ -lsec256k1 -std=c++23 secpbench.cpp && ./a.out | tail -n 6
signer thread 14 signed 210000
signer thread 16 signed 210000
signer thread 10 signed 210000
signer thread 0 signed 210000
Total signatures: 6499980
Time taken: 8859 ms
12:55:56 james@fido james/tmp %
```

# brief measurement digression


- P2P as the event loop: TX, CMPCTBLOCK, BLOCKTXN, ...
- There are basically two files in Core for perf:
  - validation.cpp
  - net\_processing.cpp
- Script caching: cached on mempool acceptance
- UTXO set cache warming

bandwidth distribution

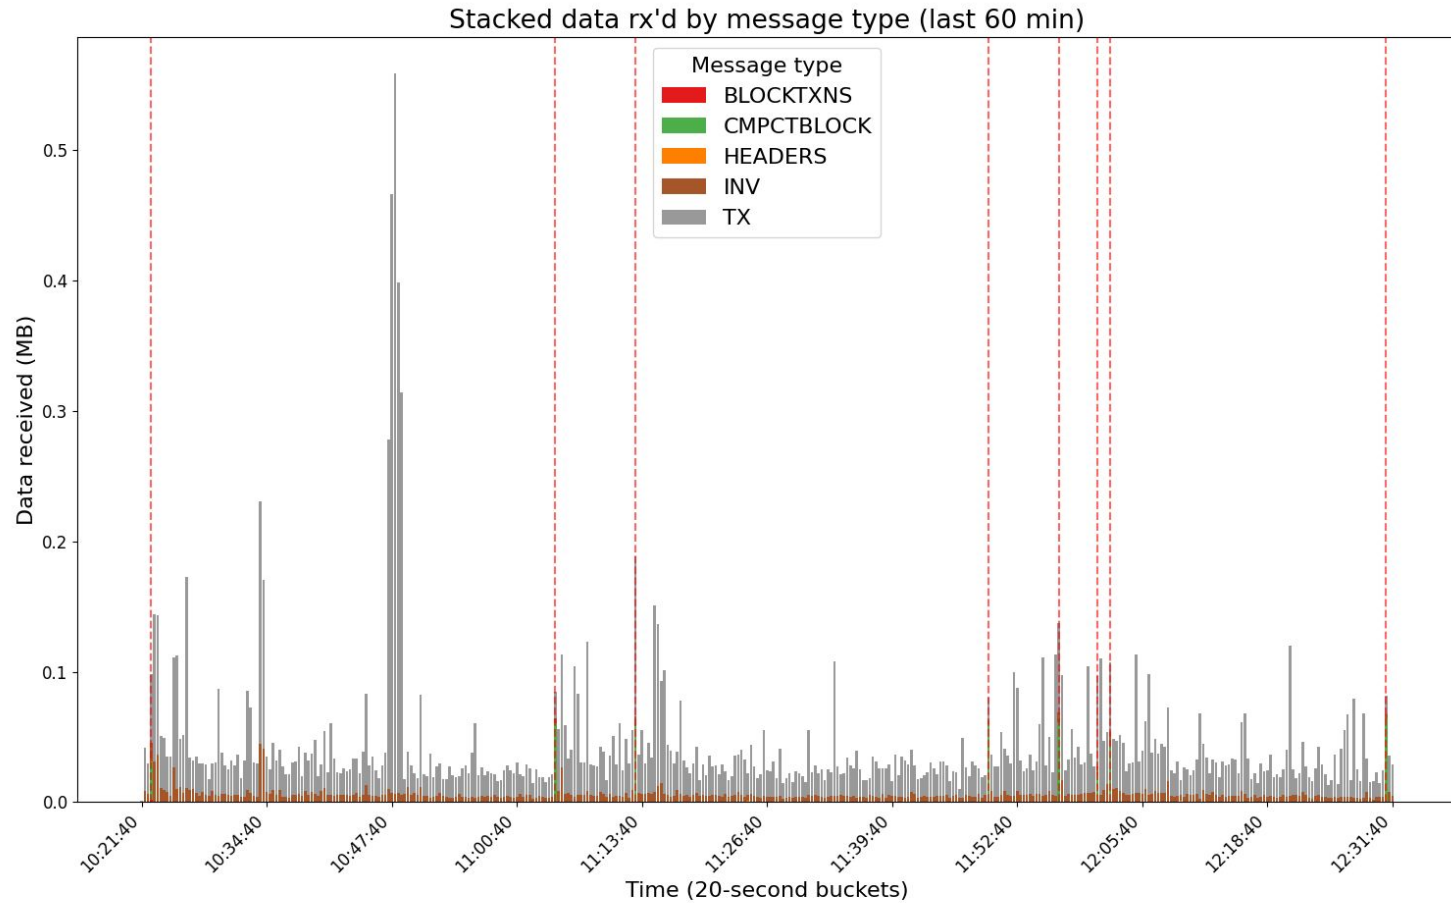





## measurement: distribution of data rx'd



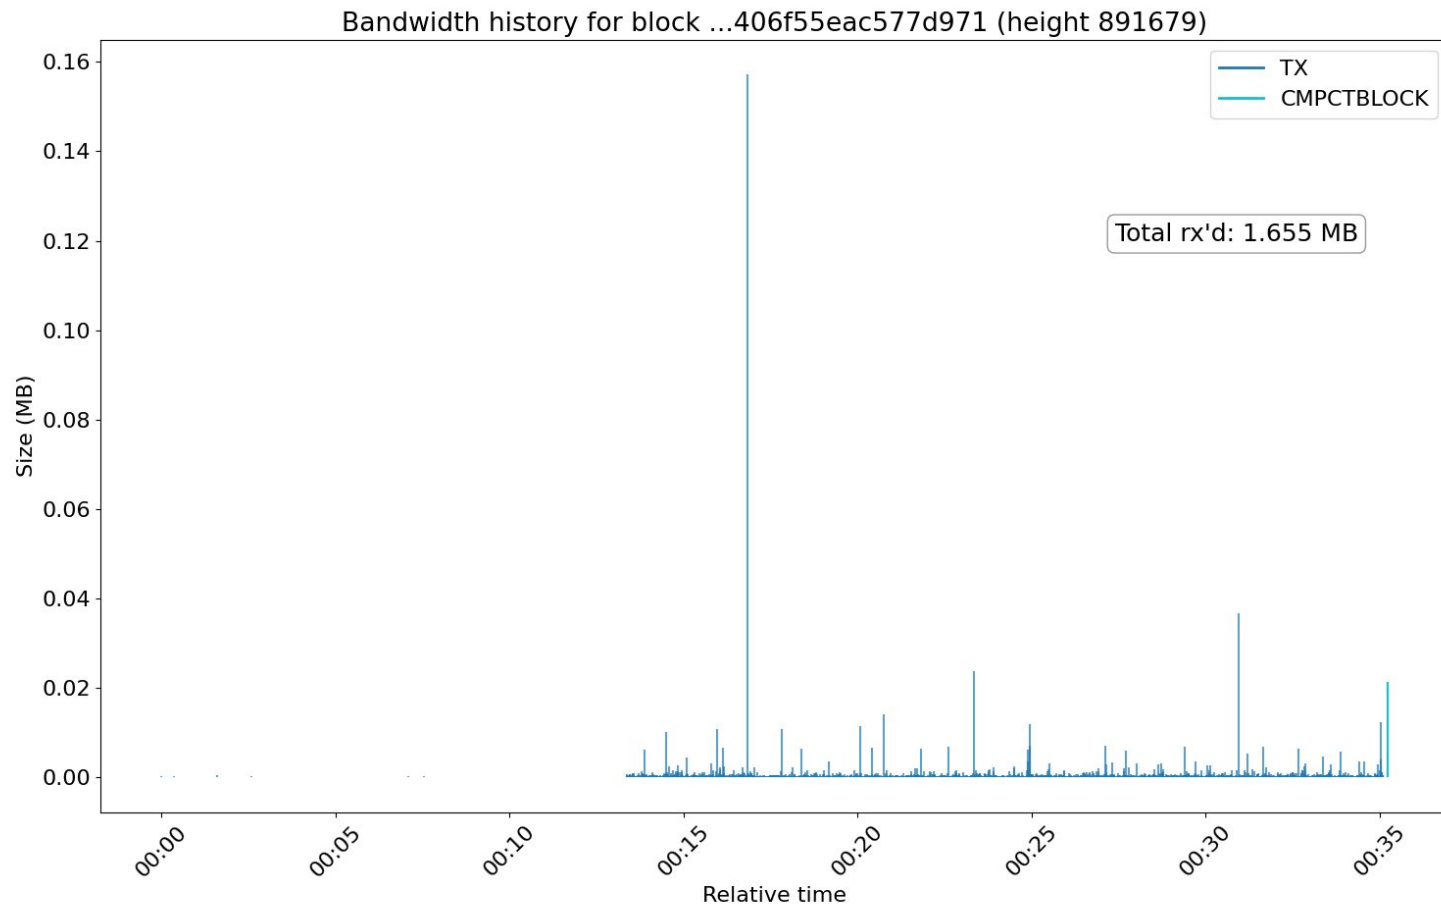
## measurement: distribution of data rx'd



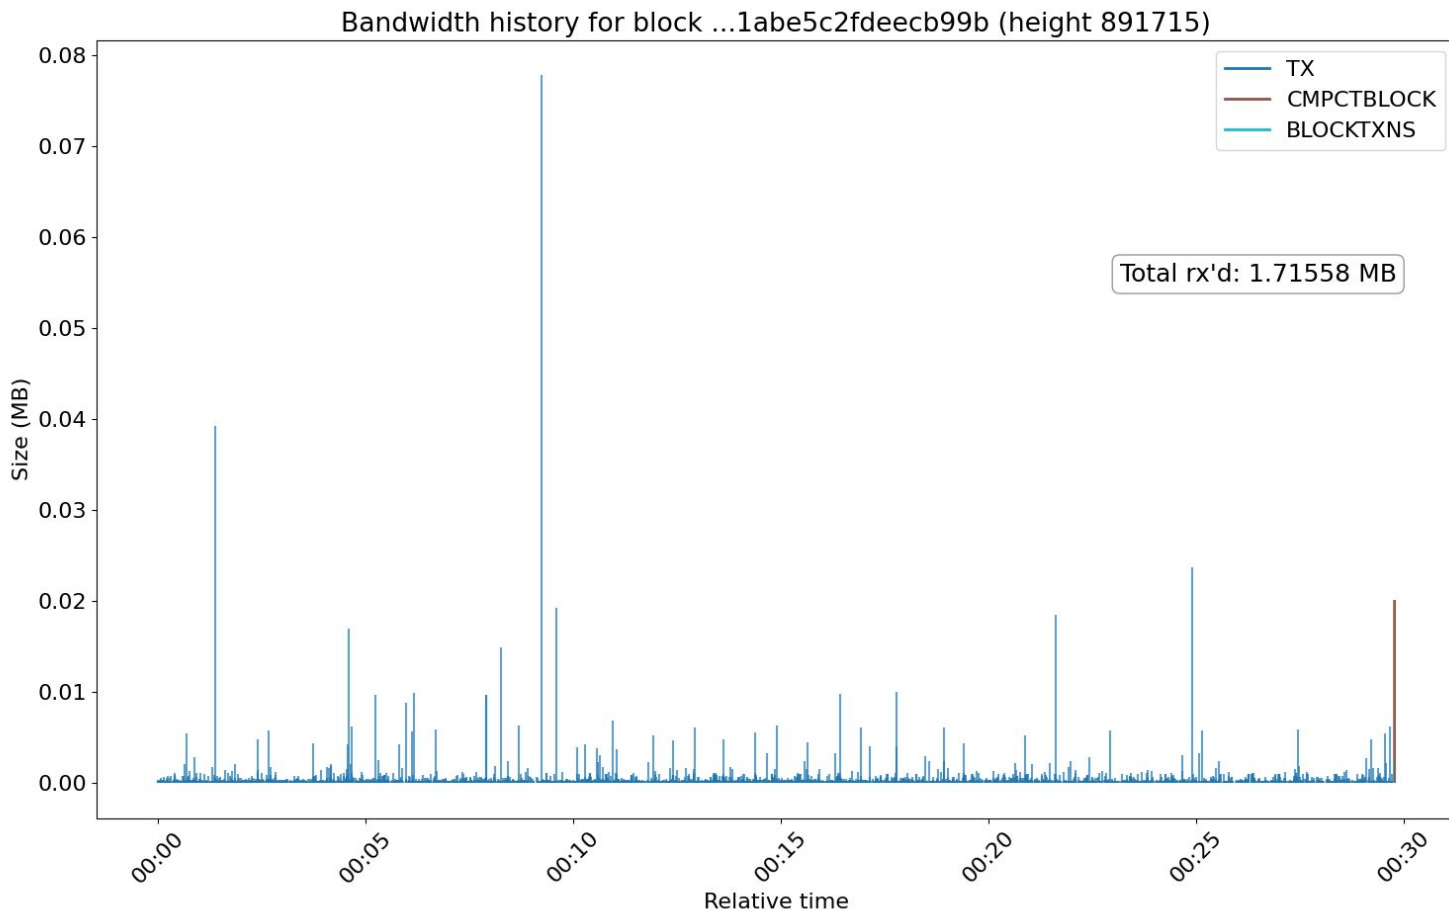
## measurement: block bandwidth history



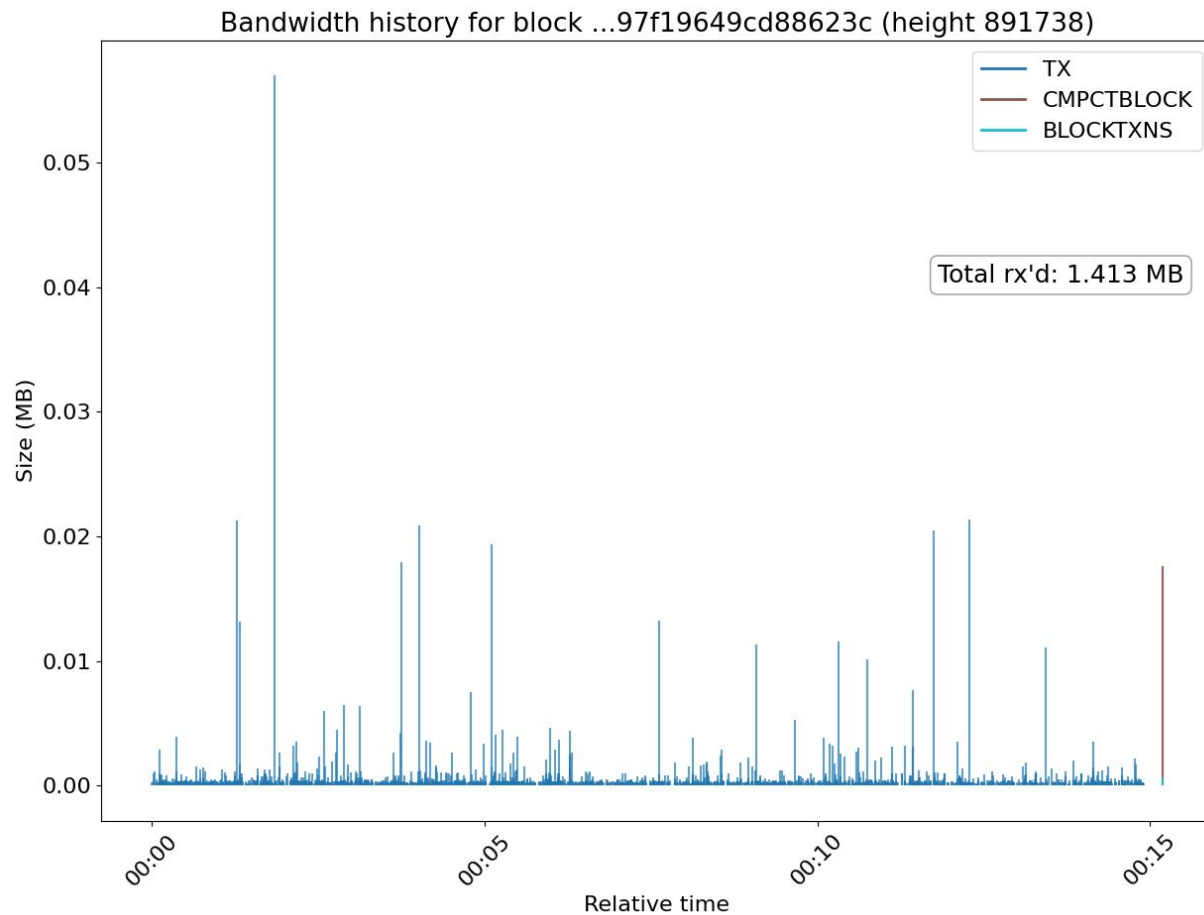
## measurement: block bandwidth history



## measurement: block bandwidth history



## measurement: block bandwidth history



# work distribution

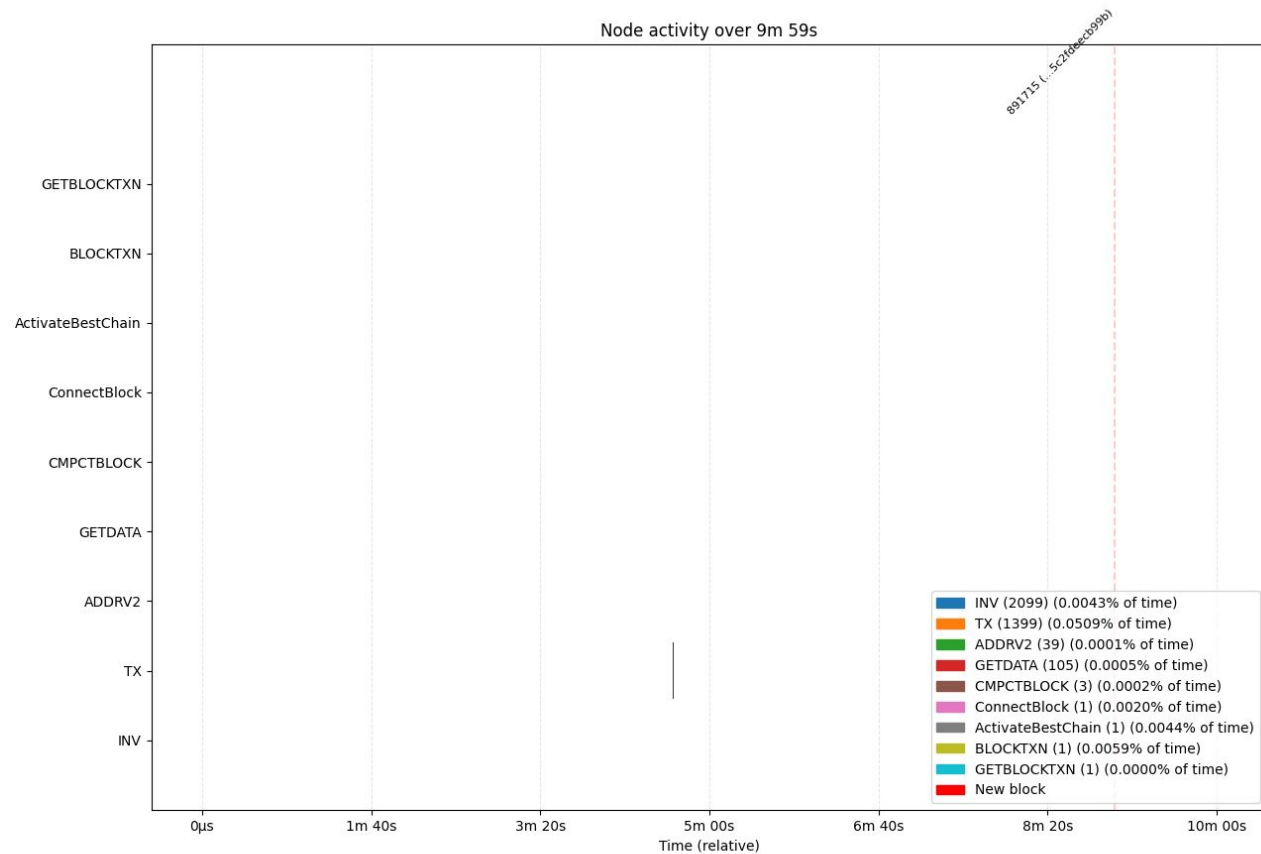
<div>**** ↓ ↑ ****</div>		@@ -2443,10 +2444,14 @@ bool Chainstate::ConnectBlock(const CBlock& block, BlockValidationState& state,	
2443	2444		AssertLockHeld(cs_main);
2444	2445		assert(pindex);
2445	2446		
2447	+	2447	+ auto event = g_event_logger->time_event("ConnectBlock");
2448	+	2448	+
2446	2449		uint256 block_hash{block.GetHash()};
2447	2450		assert(*pindex->phashBlock == block_hash);
2448	2451		const bool parallel_script_checks{m_chainman.GetCheckQueue().HasThreads()};
2449	2452		
2453	+	2453	+ event.add_metadata("blockhash="+block_hash.ToString());
2454	+	2454	+
2450	2455		const auto time_start{SteadyClock::now()};
2451	2456		const CChainParams& params{m_chainman.GetParams()};
2452	2457		
<div>**** ↓ ↑ ****</div>		@@ -2612,6 +2617,16 @@ bool Chainstate::ConnectBlock(const CBlock& block, BlockValidationState& state,	
2612	2617		}
2613	2618		}
2614	2619		
2620	+	2620	+ std::string txids{};
2621	+	2621	+
2622	+	2622	+ for (const auto& tx : block.vtx) {
2623	+	2623	+ if (txids.size() > 0) {
2624	+	2624	+ txids += ",";
2625	+	2625	}
2626	+	2626	+ txids += tx->GetHash().ToString();
2627	+	2627	}
2628	+	2628	+ event.add_metadata("txids=" + txids);
2629	+	2629	+
2615	2630		// Enforce BIP68 (sequence locks)

work distribution





# measurement: critical-path work timing



Note: 3521 events shorter than 356μs have been enlarged for visibility

## measurement: critical-path work timing

Total time: 9h 20m 47s

event type	count	total time	
TX	133684	30.537s total	0.0908% of total
ActivateBestChain	54	3.934s total	0.0117% of total
BLOCKTXN	41	3.599s total	0.0107% of total
INV	147322	1.545s total	0.0046% of total
CMPCTBLOCK	152	1.001s total	0.0030% of total
ConnectBlock	53	571ms total	0.0017% of total
GD	10946	329ms total	0.0010% of total
ADDRV2	2576	72ms total	0.0002% of total
VERSION	80	9ms total	0.0000% of total
HEADERS	212	4ms total	0.0000% of total
VERACK	62	3ms total	0.0000% of total
GETHEADERS	52	1ms total	0.0000% of total
GETBLOCKTXN	4	106µs total	0.0000% of total
ADDR	1	5µs total	0.0000% of total

## measurement: critical-path work timing

Total time: 12h 13m 12s

event type	count	total time	
TX	174830	40.577s total	0.0922% of total
ActivateBestChain	79	4.123s total	0.0094% of total
BLOCKTXN	43	3.307s total	0.0075% of total
INV	219037	2.047s total	0.0047% of total
CMPCTBLOCK	184	1.590s total	0.0036% of total
ConnectBlock	78	612ms total	0.0014% of total
GD	7166	193ms total	0.0004% of total
FlushStateToDisk	1	168ms total	0.0004% of total
HEADERS	380	91ms total	0.0002% of total
ADDRV2	3582	84ms total	0.0002% of total
GETHEADERS	70	9ms total	0.0000% of total
VERSION	102	7ms total	0.0000% of total
VERACK	68	3ms total	0.0000% of total
ADDR	1	1µs total	0.0000% of total

-----  
Total time spent working: 52.816s

Total time spent working (percent): 0.1201%

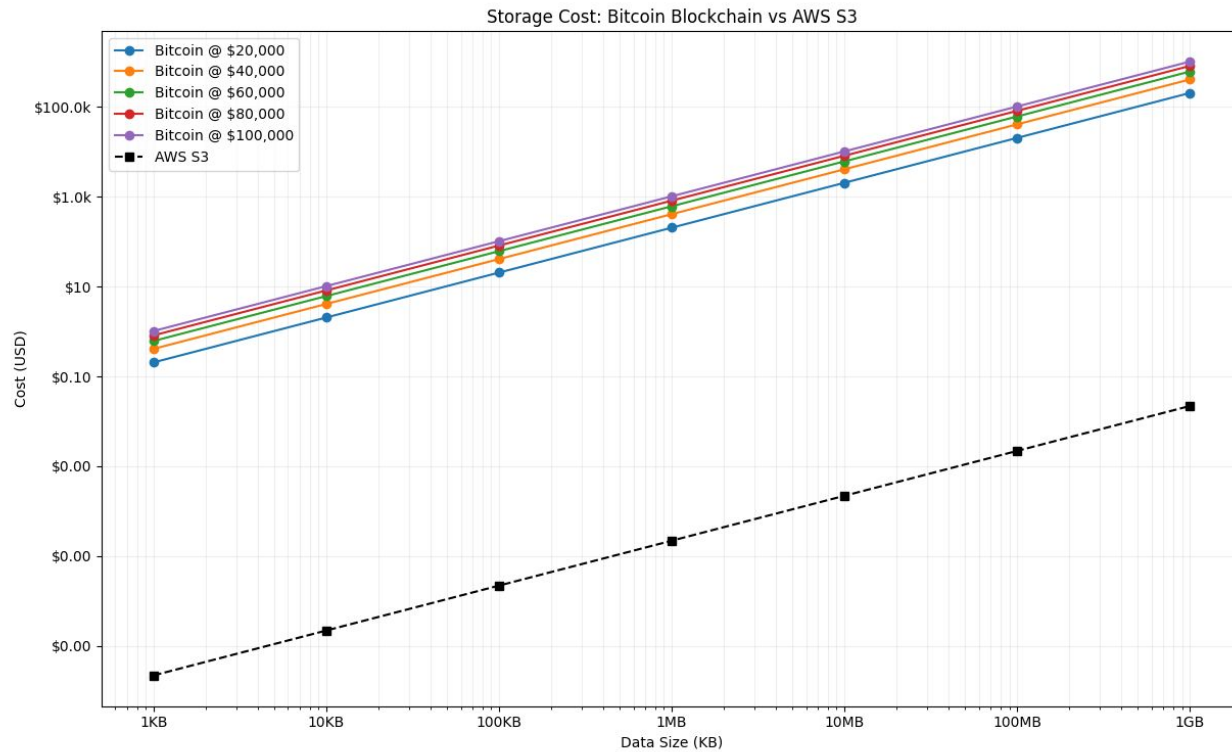
Scaled linearly 100x: 12.0%

Scaled linearly 1000x: 120.1%

# bandwidth required

Definition	Recommended up/down	Total MB/min	Total MB/hr
<b>480p/Standard</b>	1.0 Mbps/600 Kbps	13.5	810
<b>720p</b>	2.6/1.8 Mbps	22.5	1350 (1.35GB)
<b>1080p</b>	3.8/3.0 Mbps	45	2700 (2.7GB)

# spam on-chain



actual problems

```
[46]: # From gettxoutsetinfo
txouts_on_disk = 188194369
txout_size_on_disk_bytes = 12777162109
avg_utxo_size_bytes = txout_size_on_disk_bytes / txouts_on_disk

ppl = 7_000_000_000

print(f"average utxo size on disk: {avg_utxo_size_bytes:.2f} bytes")

total_utxo_size_gb = (avg_utxo_size_bytes * ppl * 2) / (1000 ** 3)

print(f"total UTXO set size (2 utxo per person): {total_utxo_size_gb:.1f}GB")

average utxo size on disk: 67.89 bytes
total UTXO set size (2 utxo per person): 950.5GB
```

```

4 //**
5 /**
6  * Benchmark for CCoinsViewCache retrieval performance.
7  *
8  * This benchmark tests the performance of retrieving coins from a CCoinsViewCache
9  * that sits in front of a LevelDB-backed CCoinsViewDB. It creates a large number
10  * of coins, stores them in the database, and then measures the time it takes to
11  * perform random retrievals.
12  *
13  * The benchmark mimics real-world usage by:
14  * 1. Setting up a CCoinsViewDB using LevelDB in a temporary directory
15  * 2. Creating a CCoinsViewCache with a 10GB maximum size
16  * 3. Populating the database with coins
17  * 4. Performing random retrievals to measure performance
18  */
19
20 #include "dbwrapper.h"
21 #include <bench/bench.h>
22 #include <coins.h>
23 #include <txdb.h>
24 #include <validation.h>
25 #include <util/fs.h>
26 #include <util/time.h>
27 #include <primitives/transaction.h>
28 #include <random.h>
29 #include <script/script.h>
30
31 #include <chrono>
32 #include <iostream>
33 #include <vector>
34 #include <random>
35
36 static const uint64_t NUM_COINS_TO_CREATE = 6'500'000'000ULL;
37 static const uint64_t NUM_RETRIEVALS = 6'500'000ULL;
38 static const uint64_t MAX_CACHE_SIZE_BYTES = 10ULL * 1024 * 1024 * 1024; // 10GB
39
40 static void CoinsViewCacheBenchmark(benchmark::Bench& bench)
41 {
42     fs::path tempDir = fs::temp_directory_path() / "bitcoin_bench_coinsview";
43     fs::create_directories(tempDir);
44
45     fs::remove_all(tempDir);
46     fs::create_directories(tempDir);
47
48     auto dbparams = DBParams{tempDir, 1 << 24};
49     auto coins_opts = CoinsViewOptions{};
50

```




# UTXO set indexing/size

- maybe ditch leveldb?
- (remember, UTXO retrieval is distributed uniformly throughout time in the average case)
  - coins cache gets warm
- Utreexo fixes this at the cost of ~2x bandwidth

mitigations

design  
for exit

## Congestion Controlled Transactions



“Did they get you to trade  
your heroes for ghosts?

Cold comfort for change?

Did you exchange  
a walk-on part in the war  
for arm/v7 as a target architecture?”

Roger Waters

Existing tech is underutilized:

cmpctblocks, assumeutxo, pruning

Prospective tech:

utreexo, weakblocks, not-leveldb

thank you

<https://github.com/jamesob/blocksize-talk>